

# CS 5523: Operating Systems

## Homework 2

**!!!! No late HW will be accepted !!!!!**

(Check BlackBoard Learn for due date and submission)

### Objective

- Learn and practice CPU scheduling algorithms by implementing them
- Learn and practice process/thread synchronization mechanisms by implementing them
- Practice system calls and library functions

### Description

In this homework, you are asked to implement a multithreaded program that will allow us to measure the performance (*i.e.*, CPU utilization, Throughput, Turnaround time, Waiting time, and Response time) of the four basic CPU scheduling algorithms (*namely*, FIFO, SJF, PR, and RR) by simulating the processes whose priority, sequence of CPU burst time(ms) and I/O burst time(ms) will be given in an input file.

*Assume that all scheduling algorithms except RR will be non-preemptive, and all scheduling algorithms except PR will ignore process priorities (i.e., all processes have the same priority in FIFO, SJF and RR). Also assume that there is only one IO device and all IO requests will be served using that device in a FIFO manner.*

---

Your program will take the name of the scheduling algorithm, related parameters (if any), and an input file name from command line. Here how your program should be executed:

```
prog -alg [FIFO|SJF|PR|RR] [-quantum [integer(ms)]] -input [file name]
```

The output of your program will be as follows:

```
Input File Name      : file name
CPU Scheduling Alg   : FIFO|SJF|PR|RR (quantum)
CPU utilization      : ....
Throughput           : ....
Turnaround time      : ....
Waiting time         : ....
Response time        : ....
```

---

The input file is formatted such that each line starts with `proc`, `sleep`, `stop` key words. Following `proc`, there will be a sequence of integer numbers: the first one represents the priority (1: lowest, ..., 5: normal, ..., 10: highest), while the remaining ones represent CPU burst and I/O burst times (ms) in an alternating manner. Following `sleep`, there will be an integer number representing the time (ms) after which there will be another process. So one of the threads in your program would be responsible for processing this file as follows. As long as it reads `proc`, it will create a new process and put it in a ready queue (clearly this process is not an actual one, it will be just a simple data structure (similar to PCB) that contains the given priority and the sequence of CPU burst and I/O burst times, and other fields). When this thread reads `sleep x`, it will sleep `x` ms and then try to read new processes from the file. Upon reading `stop`, this thread will quit.

---

Here is a sample input file:

---

```
proc  1 10 20 10 50 20 40 10
proc  1 50 10 30 20 40
sleep 50
proc  2 20 50 20
stop
```

---

You need at least two other threads to simulate the behaviors of CPU scheduler and I/O system.

- CPU scheduler thread will check ready queue; if there is a process, it will pick one according to the scheduling algorithm from ready queue and hold CPU resource for the given CPU burst time (or for quantum time if the scheduling algorithm is RR). Then it will release CPU resource and put this process into IO queue (or ready queue if RR is used) or just terminate if there is no more CPU burst. Then CPU scheduler thread will check ready queue again and repeat the same...
  - I/O system thread will check IO queue; if there is a process, it will hold IO device for the given IO burst time and then put this process back into ready queue. Then it will check IO queue and repeat the same ....
- 

Basically you will have at least the above mentioned three threads and the main one. These threads need to be synchronized as they cooperate to collect data for performance measures and share data through ready and IO queues. Main thread will wait until the file reading thread is done and the ready queue and IO queue are empty, then it will print the performance evaluation results and terminate the program...

You are free to design and implement this program in C/C++ or Java (but if you use Java, DO NOT use the ‘*synchronized*’ methods or high-level thread-safe Java classes. Instead use some form of semaphore or basic synchronization mechanisms in Java. Also `synchronized(obj) { }` structure is OK to use for protecting critical sections in your methods).

Also you can use any data structures in different parts of your program and share them along with new variables; but when it comes to maintaining the list of processes in ready queue or IO queue, we would like you to use **double linked list**, as this is the case in many practical OS.

---

**Grading:** This is a 200-point homework.

First make sure your program works for FIFO CPU scheduling. This will be **100** points.

Then you can add/test the others: SJF (**25** points), PR (**25** points) and RR (**40** points).

Write a 2-3 page **report** (**10** points) to describe your design choices at the high level and your results.

Do all your work under a directory **lastname\_hw2**, which should include your source codes, instructions to compile/execute, and some output files showing your test results etc...

Zip **lastname\_hw2** and submit **lastname\_hw2.zip** through BB Learn

---

## Submission

You must submit your work using Blackboard Learn and respect the following rules:

- 1) All assignments must be submitted as either a zip or tar archive file unless it is a single pdf file.

- 2) Assignments must include all source code.
- 3) Assignments must include an output.txt file which demonstrates the final test output run by the student.
- 4) If your assignment does not run/compile, the output.txt file should include an explanation of what was accomplished, what the error message was that prevented the student from finishing the assignment and what the student BELIEVES to be the underlying cause of the error.